

Message 21 · Architect · 12:15

I would correct G2 to No. [...]

For G3, the model decides at runtime which registered internal source is relevant to each requirement and may combine evidence from several sources. [...]

The route is therefore Agent, based on adaptive source selection, query reformulation and stopping decisions.

Message 24 · Assistant · 12:18

Your answer settles G3: at most two internal systems and two cross-system fields, with the web closing a separate gap and not part of the join. G3 does **not fire**. [...]

Route: agent, fixed. Your priorities can't override it. [...]

M on P2 creates a conflict on this route. The agent's safe-failure capability is the lowest, so M cannot be met by design alone. It would be a deployment conflict, resolved by re-scoping, moving the requirement outside the task, human closure of the loop, or not automating. [...]

Message 22 · Assistant · 12:15

I accept G2 as **does not fire**. [...]

I can't close G3 yet. The gates run in a fixed order and G3 comes before G4 and G5, so I record your G4 and G5 answers (both fire) but they don't set the route while G3 is open. [...]

In a demanding but realistic case, how many separate internal systems (document store, contact records, knowledge base, and so on) and how many fields would the model have to combine into a single answer? [...]

Message 25 · Architect · 12:42

[...] My performance weights are: [...]

- P2 Safe failure: M. [...]

I would resolve the P2 conflict by moving the final readiness authority outside the agent loop. The agent may propose *ready*, but deterministic post-checks must verify that every required field has attributable evidence, all mandatory sources were checked, citations resolve, freshness rules pass, and no conflict, placeholder, incomplete read or failed tool call remains. [...]

Message 23 · Architect · 12:17

In a demanding but still normal case, the model would combine information from at most two separate internal systems—typically a contact or customer record plus one document or knowledge source—and resolve no more than two cross-system fields into the context package. [...]

Message 26 · Assistant · 12:43

Build Context gathering as a guarded ReAct agent loop. [...]

P2 resolution. The rulebook's own row for this case is post-state validation, with human approval only before consequential actions. There is no consequential action here, so no approval gate is needed. Your deterministic readiness authority outside the loop is that post-state validation, and I accept it. **Residual risk, recorded:** deterministic checks can verify only what is checkable (evidence present, sources checked, citations resolve, freshness, no open read or failed tool call). The agent's measured failure is that wrong outputs look valid. So a requirement that was never derived, or a relevant source that was never selected, can still pass unnoticed. [...]