

Claude · message 10 · selected parts of the verdict

Verdict: build an LLM workflow.

Code owns the flow: it reads the sources, makes one bounded model step per entry to write the line items plus one for the intro and closing text, then fills the template and saves the draft. No gate fired, and on your weights the workflow leads 37 to 32, mainly on safe failure, traceability and schema guarantee, which you rated critical.

Mandatory minimum for this workflow: one bounded re-plan round for empty reads, a placeholder guard before dispatch (no unresolved value reaches the template), per-item calls with a deterministic count check, fail-closed conditions, and a dev-slice pilot before the build is frozen. [...]

Fit and profile record

Criterion	Applies	Weight	Workflow cap.	Agent cap.	Workflow	Agent
T1 Rule application	No	0	3	2	0	0
P1 Traceability	Yes	3	3	2	9	6
P2 Safe failure	Yes	3	3	1	9	3
P3 Schema guarantee	Yes	3	3	2	9	6
P4 Efficiency	Yes	1	3	2	3	2
P5 Messy input	Yes	2	2	3	4	6
P6 Evolvability	Yes	3	1	3	3	9
Total					37	32

[...] Margin: A margin of five is not a close call. The P6 scores are derived, not measured.

Components and constraints

P5 (workflow capability 2, weight 2): An input-normalisation stage is mandatory.

P6 (workflow capability 1, weight 3): This is a residual risk. Structural changes, such as new template fields, a different client-file layout or different Notion columns, require explicit review and graph modification. [...]

Design knowledge

[...] 02 Build form. [...]

canonical two-stage is the leading candidate. It has a deterministic retrieval core, where code reads the month's entries, client data and template with fixed queries and checks coverage through pagination, deduplication and validation. The model step then runs per entry, following the per-item map pattern. Its benefit for item-specific text is derived, not measured. [...]

Executor commitment: Will the build implement the mandatory minimum, including the dev-slice pilot?