

STAGE 1 — QUALIFICATION & AGENT-GATES (top to bottom; the first 'yes' ends the pass)

#	Category	Question and how to answer it	Answer	If YES →
Q1	LLM necessity	Can the required output be produced using only predefined deterministic logic? If no semantic interpretation, context-dependent reasoning, or context-specific language generation is required.	no	Implement deterministically
Q2	Atomic responsibility	Does the subtask combine multiple executable and testable LLM responsibilities? Examples: researching and approving · extracting facts and applying rules · deciding and executing an irreversible action. Atomicity refers to ONE coherent goal and success criterion, not to the number of LLM calls.	no	Decompose
Q3	Granularity	Does the responsibility exceed the tested boundary of one reliable task execution? Warning conditions include: six or more independent output requirements · semantic combination across three or more systems · a separate retrieval or reasoning process for each item.	no	Decompose
Q4	Adaptive control	Must the LLM select or create a next step that cannot be fully predefined? If runtime observations require an OPEN choice of the next tool, action, intermediate step, or stopping point — including unexpected outcomes, outcome-dependent repetition whose stop condition CANNOT be checked deterministically, and the observable state changes that require replanning.	no	Agent loop
Q5	Adaptive search	Does the subtask require adaptive search over an open or unknown search space? If the next query or source depends on previous findings, the relevant source is unknown, or completeness cannot be verified deterministically in a semantically ambiguous search space.	no	Agent loop

STAGE 2 — WEIGHTED PARADIGM FIT (only if all gates are 'no')

#	Category	Criterion and how to answer it	Applies	Weight	eff. w	LLM workflow	Agent loop
T1	Rule success	Can the goal be reached by applying clear written rules? If the rules can be retrieved at runtime (policy store, thresholds, escalation matrix). Vague, conflicting, or discretionary rules do not qualify.	yes	3	3	3	2
T2	Multi-item extraction success	Must different item-specific values be completely extracted from unstructured text and matched to the correct items? Applies when one natural-language input contains multiple item-value mappings. Identical values or already structured mappings do not qualify.	yes	3	3	2	3
T3	Retrieval success	Can the required information be retrieved from known and enumerable sources? If all relevant sources are known and coverage can be checked through predefined querying, pagination, deduplication, and validation.	no	3	0	3	1
P1	Efficiency	Efficiency & predictability Cost, latency, and budget predictability per solved task.	•	3	3	3	1
P2	Governance	Traceability & audit evidence Ability to reconstruct tool calls, retrieved evidence, state changes, and outputs.	•	2	2	3	2
P3	Governance	Safe failure direction Ability to constrain errors and uncertainty toward a safe state (e.g. human escalation).	•	3	3	3	1
P4	Output control	Output schema guarantee Ability to produce typed and validated output (fixed, machine-validated format).	•	2	2	3	2
P5	Robustness	Robustness to messy input Tolerance to noisy, informal, incomplete, or imperfectly structured input.	•	3	3	2	3
P6	Evolvability	Flexibility under change Ability to accommodate changing task definitions or input distributions.	•	0	0	1	3

LLM workflow						51	89%
Agent loop						38	67%

S = Σ applies × weight × capability (deployment criteria always apply; M counts as weight 3 AND requires capability ≥ 2 , non-compensable).